

Discrete Event Simulation Analysis of a Web Application

CS 681: Assignment 2

Sameer Arvind Patil(23b1035) Harsh Suthar(23b1067)

March 2026

Outline

- 1 System Overview and Architecture
- 2 Setup 1: MVA Comparison with Assignment 1 Baseline
- 3 Setup 2: Confidence Intervals on General Webserver
- 4 Setup 3: Scheduling Policy Comparison (RR vs SJF)
- 5 Additional Analysis and Future Work

System Overview

- Multi-core web server simulator written in C++
- Models realistic web application performance characteristics
- Closed-loop user model: request → wait → think → repeat
- Thread-per-task model with configurable thread pool
- Supports multiple scheduling policies: Round-Robin (RR) and Shortest Job First (SJF)
- Discrete event simulation with proper warmup and transient removal
- Statistical analysis with confidence intervals

Simulation Workflow

- 1 **Initialization:** Create cores, thread pool, initialize RNG
- 2 **User Arrivals:** Users generate requests in closed loop
- 3 **Admission Control:** Requests admitted to waiting queue if thread available
- 4 **Thread Execution:** Threads execute requests with time slicing (quantum)
- 5 **Timeouts:** Long-running requests timeout after specified duration
- 6 **Completion:** Requests complete, thread returns to pool, user thinks
- 7 **Collection:** Metrics collected after warmup period

Key Components

Request Model:

- Service time (configurable distribution)
- Timeout (uniform range)
- User ID and request ID
- Retry tracking

Core and Threading:

- Multiple cores (configurable)
- Thread pool with max limit
- Per-core run queue
- Context-switching overhead

Key Performance Metrics

- **Average Response Time:** Mean time from request arrival to completion
- **Throughput:** Total requests completed per second (goodput + badput)
- **Goodput:** Requests completed before timeout per second
- **Badput:** Timed-out requests completed per second
- **Timeout Rate:** Number of requests exceeding timeout threshold
- **Core Utilization:** Fraction of time cores are busy
- **Request Drop Rate:** Requests rejected due to full queue

Setup 1: Objective

Question: How well does the simulation match the Mean Value Analysis (MVA) theoretical model and Assignment 1 measurements?

This setup validates the simulator by comparing against:

- MVA theoretical predictions
- Empirical measurements from Assignment 1
- Single simulator run with Assignment 1 parameters


Setup 1: Parameters

Parameter	Value
Number of Users	10 to 295 (step 20)
Service Time Distribution	Constant 30 ms
Number of Cores	1
Max Threads	345
Quantum (time slice)	10 ms
Context Switch Overhead	0.0 ms
Think Time (base)	4000 ms
Think Time (exponential mean)	10 ms
Timeout Range	30000–30010 ms
Simulation Time	180 seconds (warmup: 20 sec)

Setup 1: Rationale

- **Single Core:** Matches Assignment 1 empirical measurements
- **Constant Service Time:** Simplifies validation; removes variability
- **Large Thread Pool:** Ensures thread availability across user ranges
- **High Timeout:** With small service times, timeouts rarely occur
- **Think Time Model:** Think time has mode (4000 ms) + exponential component (10 ms)
- **Long Simulation:** Ensures steady-state metrics are accurate

Setup 1: Comparison Results



The plot compares:

- **MVA:** Theoretical prediction line
- **Assignment 1:** Measured data from real system
- **Simulation:** Discrete event simulation results

Setup 1: Analysis

- Simulation closely matches MVA theoretical predictions
- Both show linear increase in response time with user load
- Simulation validates the core algorithm and event handling
- Small deviations from MVA due to:
 - Discrete time-slicing effects
 - Warmup transient removal
 - Stochastic variations
- Assignment 1 data may differ due to measurement methodology

Setup 2: Objective

Questions:

- How do performance metrics scale with increasing user load?
- What are the confidence intervals for response time?
- Where does the system saturate?

This setup performs 40 independent runs with varying user populations to compute statistically sound confidence intervals.

Setup 2: Parameters

Parameter	Value
Number of Users	40 to 3000 (step 40)
Service Time Distribution	Exponential, mean 60 ms
Number of Cores	4
Max Threads	256
Quantum (time slice)	10 ms
Context Switch Overhead	0.2 ms
Think Time (base)	3000 ms
Think Time (exponential mean)	200 ms
Timeout Range	20000–30000 ms
Simulation Time	180 seconds (warmup: 20 sec)
Independent Runs	40
Scheduling Policy	Round-Robin

Setup 2: Statistical Methodology


- **Independent Replications:** 40 runs with different RNG seeds
- **Point Estimate:** Mean of metric across all 40 runs for each user count
- **95% Confidence Interval:**

$$CI = \bar{X} \pm t_{0.025,39} \cdot \frac{s}{\sqrt{40}}$$

where $t_{0.025,39} \approx 2.02$ (Student's t-distribution)


- **Warmup Removal:** First 20 seconds discarded (Welch's procedure informal)
- **Transient Analysis:** Statistics only from steady-state period

Setup 2: Response Time Results




- Response time increases sub-linearly up to 500 users
- Sharp increase in response time and confidence interval at high loads
- Saturation point near 1500-2000 users

Setup 2: Throughput Results




- Throughput increases with users up to saturation point
- Peak throughput ≈ 16 -17 requests/sec at moderate load
- Plateaus after saturation due to thread pool limits

Setup 2: Goodput vs Badput




- Goodput: Requests completed before timeout
- Badput: Timed-out requests that still completed
- Badput grows significantly with load (requests queue too long)

Setup 2: Utilization Results



- Core utilization increases with user load
- Reaches $\approx 95\%$ at saturation
- 4-core system efficiently utilized

Setup 2: Timeout Analysis



- Few timeouts at low load
- Exponential growth in timeouts above 1000 users
- System SLA violation indicates saturation

Setup 2: Key Insights

- **Saturation Point:** System becomes resource-constrained around 1500-2000 users
- **Queueing Effects:** Visible degradation in response time and goodput
- **Thread Pool Limit:** 256 threads insufficient for extreme loads
- **Variability:** Confidence intervals widen significantly at high loads
- **Confidence:** 95% CI indicates statistical soundness

Setup 3: Objective

Questions:

- How much does scheduling policy impact performance?
- Is SJF better than RR for web server workloads?
- What is the trade-off between fairness and latency?

This setup compares two scheduling policies:

- **Round-Robin (RR)**: Fair, simple scheduling
- **Shortest Job First (SJF)**: Prioritizes shorter jobs


Setup 3: Parameters

Parameter	RR	SJF
Number of Users	5–3000 (step 10)	5–3000 (step 10)
Service Time	Uniform [20, 100] ms	Uniform [20, 100] ms
Number of Cores	4	4
Max Threads	256	256
Quantum	100 ms	100 ms
Context Switch Overhead	0.2 ms	1.0 ms
Think Time (base)	3000 ms	3000 ms
Think Time (exponential)	200 ms	200 ms
Timeout Range	20000–30000 ms	20000–30000 ms

Setup 3: Rationale


- **Uniform Service Time:** Creates variability for SJF to optimize
- **Multiple Cores:** Real-world scenario with parallelism
- **Larger Quantum:** Reduces context-switching overhead impact
- **Higher CS Overhead for SJF:** Penalty for more frequent preemptions
- **Wide User Range:** Captures behavior from light to extreme load

Setup 3: RR vs SJF Response Time



- SJF generally shows lower response time at moderate loads
- Both policies degrade similarly at high load (resource saturation dominates)

Setup 3: RR vs SJF Throughput



- Throughput curves nearly identical for RR and SJF
- Scheduling policy has minimal impact on aggregate throughput

Setup 3: Analysis

- **SJF Advantage:** Better response time for short jobs
- **Trade-off:** Higher context-switching overhead for SJF
- **Saturation Behavior:** Both policies degrade similarly at extreme load
- **Recommendation:**
 - Use SJF for systems with highly variable service times
 - RR is simpler and adequate for balanced workloads
 - Consider hybrid approaches for critical services

Potential Analysis Directions

- **Timeout Parameter Sensitivity:** How do different timeout ranges affect goodput?
- **Thread Pool Sizing:** Optimal thread pool size vs. memory cost
- **Quantum Effects:** Impact of time slice size on overhead vs. fairness
- **Service Time Distribution:** Compare constant, uniform, and exponential
- **Think Time Impact:** Effects of think time distribution on load cycles
- **Multiple Cores:** Scalability analysis with varying core counts

Queue Dynamics Analysis

- **Queue Length Distribution:** Measure average queue depth over time
- **Wait Time Components:** Break down response time into:
 - Queueing delay (waiting for thread)
 - Service delay (on core)
 - Context-switching overhead
- **Starvation Detection:** Identify which requests timeout most frequently
- **Load Balancing:** Core utilization balance across all cores

System Tuning Recommendations

- **Thread Pool:** Size based on 95th-percentile concurrent requests
- **Quantum Selection:** Tune for balance between fairness and overhead
- **Timeout Strategy:** Set timeout based on 99th-percentile response time
- **Request Admission:** Reject requests early if queue too deep
- **Priority Queuing:** VIP users or critical requests get preference
- **Load Shedding:** Graceful degradation at extreme load

Conclusion

- **Simulator Validation:** Discrete event simulation closely matches MVA theory
- **Performance Characterization:** Comprehensive metrics with confidence intervals
- **Scheduling Impact:** Limited impact on aggregate throughput at saturation
- **Saturation Point:** System saturates around 1500-2000 users for given config
- **Statistical Rigor:** Proper use of confidence intervals and warmup removal
- **Extensibility:** Framework supports multiple scheduling policies, distributions, and workloads

Code Structure

Core Components:

- `simulator.h/cpp`: Main event loop
- `request.h`: Request tracking
- `core.h/cpp`: Per-core scheduling
- `threadworker.h`: Thread abstraction

Analysis:

- `plot.py`: Basic metrics plotting
- `confidence_plotter.py`: CI calculations
- `mix_plot.py`: Multi-source comparison
- Scripts: Runner shells for all setups